

Predict คนลาออกจากงาน

Absurdism

ภูมิภัทร พิณธีระ 6410301032

อนันต์โชติ ธนโชคธัญสิริ 6410301033

Problem

- พนักงานเป็นฟันเฟืองหนึ่งขององค์กร
- การลาออกของพนักงานส่งผลต่อการทำงาน, ประสิทธิภาพขององค์กร
- ฉะนั้นแต่ละองค์กรควรที่จะรับรู้ว่าปัจจัย (factor) ใดที่จะส่งผลต่อการลาออกของพนักงาน เพื่อที่จะได้เตรียมการในการแก้ไขปัญหา

Goal

- ใช้เพื่อวิเคราะห์ว่าปัจจัยใดที่ส่งผลต่อการลาออกของพนักงาน

Data Set

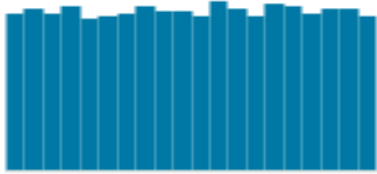
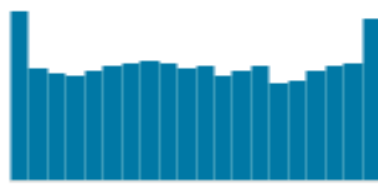
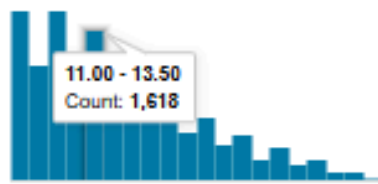

Employee Attrition Classification Dataset

Data Card

Code (30)

Discussion (1)

Suggestions (1)

Employee ID	# Age	Gender	# Years at Company	Job Role	# Months
ID	age	gender	Years at Company	Job Role	Month
		<div>Male54%</div> <div>Female46%</div>		<div>Technology26%</div> <div>Healthcare23%</div> <div>Other (7653)51%</div>	
5	74.5k	18	59	1	51
52685	36	Male	13	Healthcare	8029
30585	35	Male	7	Education	4563
54656	50	Male	7	Education	5583
33442	58	Male	44	Media	5525
15667	39	Male	24	Education	4604
3496	45	Female	30	Healthcare	8104
46775	22	Female	5	Healthcare	8700

<https://www.kaggle.com/datasets/stealthtechnologies/employee-attrition-dataset/data>

Data Understanding

```
RangeIndex: 74498 entries, 0 to 74497
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	74498 non-null	int64
1	Gender	74498 non-null	object
2	Years at Company	74498 non-null	int64
3	Job Role	74498 non-null	object
4	Monthly Income	74498 non-null	int64
5	Work-Life Balance	74498 non-null	object
6	Job Satisfaction	74498 non-null	object
7	Performance Rating	74498 non-null	object
8	Number of Promotions	74498 non-null	int64
9	Overtime	74498 non-null	object
10	Distance from Home	74498 non-null	int64
11	Education Level	74498 non-null	object
12	Marital Status	74498 non-null	object
13	Number of Dependents	74498 non-null	int64
14	Job Level	74498 non-null	object
15	Company Size	74498 non-null	object
16	Company Tenure	74498 non-null	int64
17	Remote Work	74498 non-null	object
18	Leadership Opportunities	74498 non-null	object
19	Innovation Opportunities	74498 non-null	object
20	Company Reputation	74498 non-null	object
21	Employee Recognition	74498 non-null	object
22	Attrition	74498 non-null	object

```
dtypes: int64(7), object(16)
```

```
data['Attrition'].value_counts()
```

```
[7]
```



```
0.0s
```

```
...
```

```
Attrition
```

```
Stayed      39128
```

```
Left        35370
```

```
Name: count, dtype: int64
```

```
data['Age'].describe()
```

✓ 0.0s

```
count    74498.000000
mean      38.529746
std       12.083456
min       18.000000
25%       28.000000
50%       39.000000
75%       49.000000
max       59.000000
Name: Age, dtype: float64
```

```
data['Gender'].value_counts()
```

✓ 0.0s

```
Gender
Male    40826
Female  33672
Name: count, dtype: int64
```



```
data['Years at Company'].describe()
```

```
[10] ✓ 0.0s
```

```
... count      74498.000000  
    mean         15.721603  
    std          11.223744  
    min           1.000000  
    25%           7.000000  
    50%          13.000000  
    75%          23.000000  
    max          51.000000  
    Name: Years at Company, dtype: float64
```

```
data['Years at Company'].value_counts()
```

```
[11] ✓ 0.0s
```

```
... Years at Company  
5      3084  
1      3056  
2      3039  
8      3015  
10     2987  
9      2965  
3      2961  
6      2952  
7      2933  
4      2903  
11     2893  
12     2758  
13     2547  
14     2349  
15     2281  
16     2145  
17     1994  
18     1889  
19     1754  
20     1729  
21     1606  
23     1557  
22     1537  
24     1385
```

```
data['Job Role'].value_counts()
```

```
2]
```

✓ 0.0s

Job Role

Technology 19322

Healthcare 17074

Education 15658

Media 11996

Finance 10448

Name: count, dtype: int64

Feature Engineering

```
def classify_birth_year_group(gen):
```

```
    birth_year = 2024 - gen
```

```
    if birth_year >= 2013:
```

```
        return 'Gen_Alpha'
```

```
    elif 1995 <= birth_year <= 2012:
```

```
        return 'Gen_Z'
```

```
    elif 1980 <= birth_year <= 1994:
```

```
        return 'Gen_Y'
```

```
    elif 1965 <= birth_year <= 1979:
```

```
        return 'Gen_X'
```

```
    else:
```

```
        return 'Baby_Boomer'
```

```
data['Generation'] = data['Age'].apply(classify_birth_year_group)
```

```
data
```

```
# X11 = Prepro_Data[['Generation']]
```

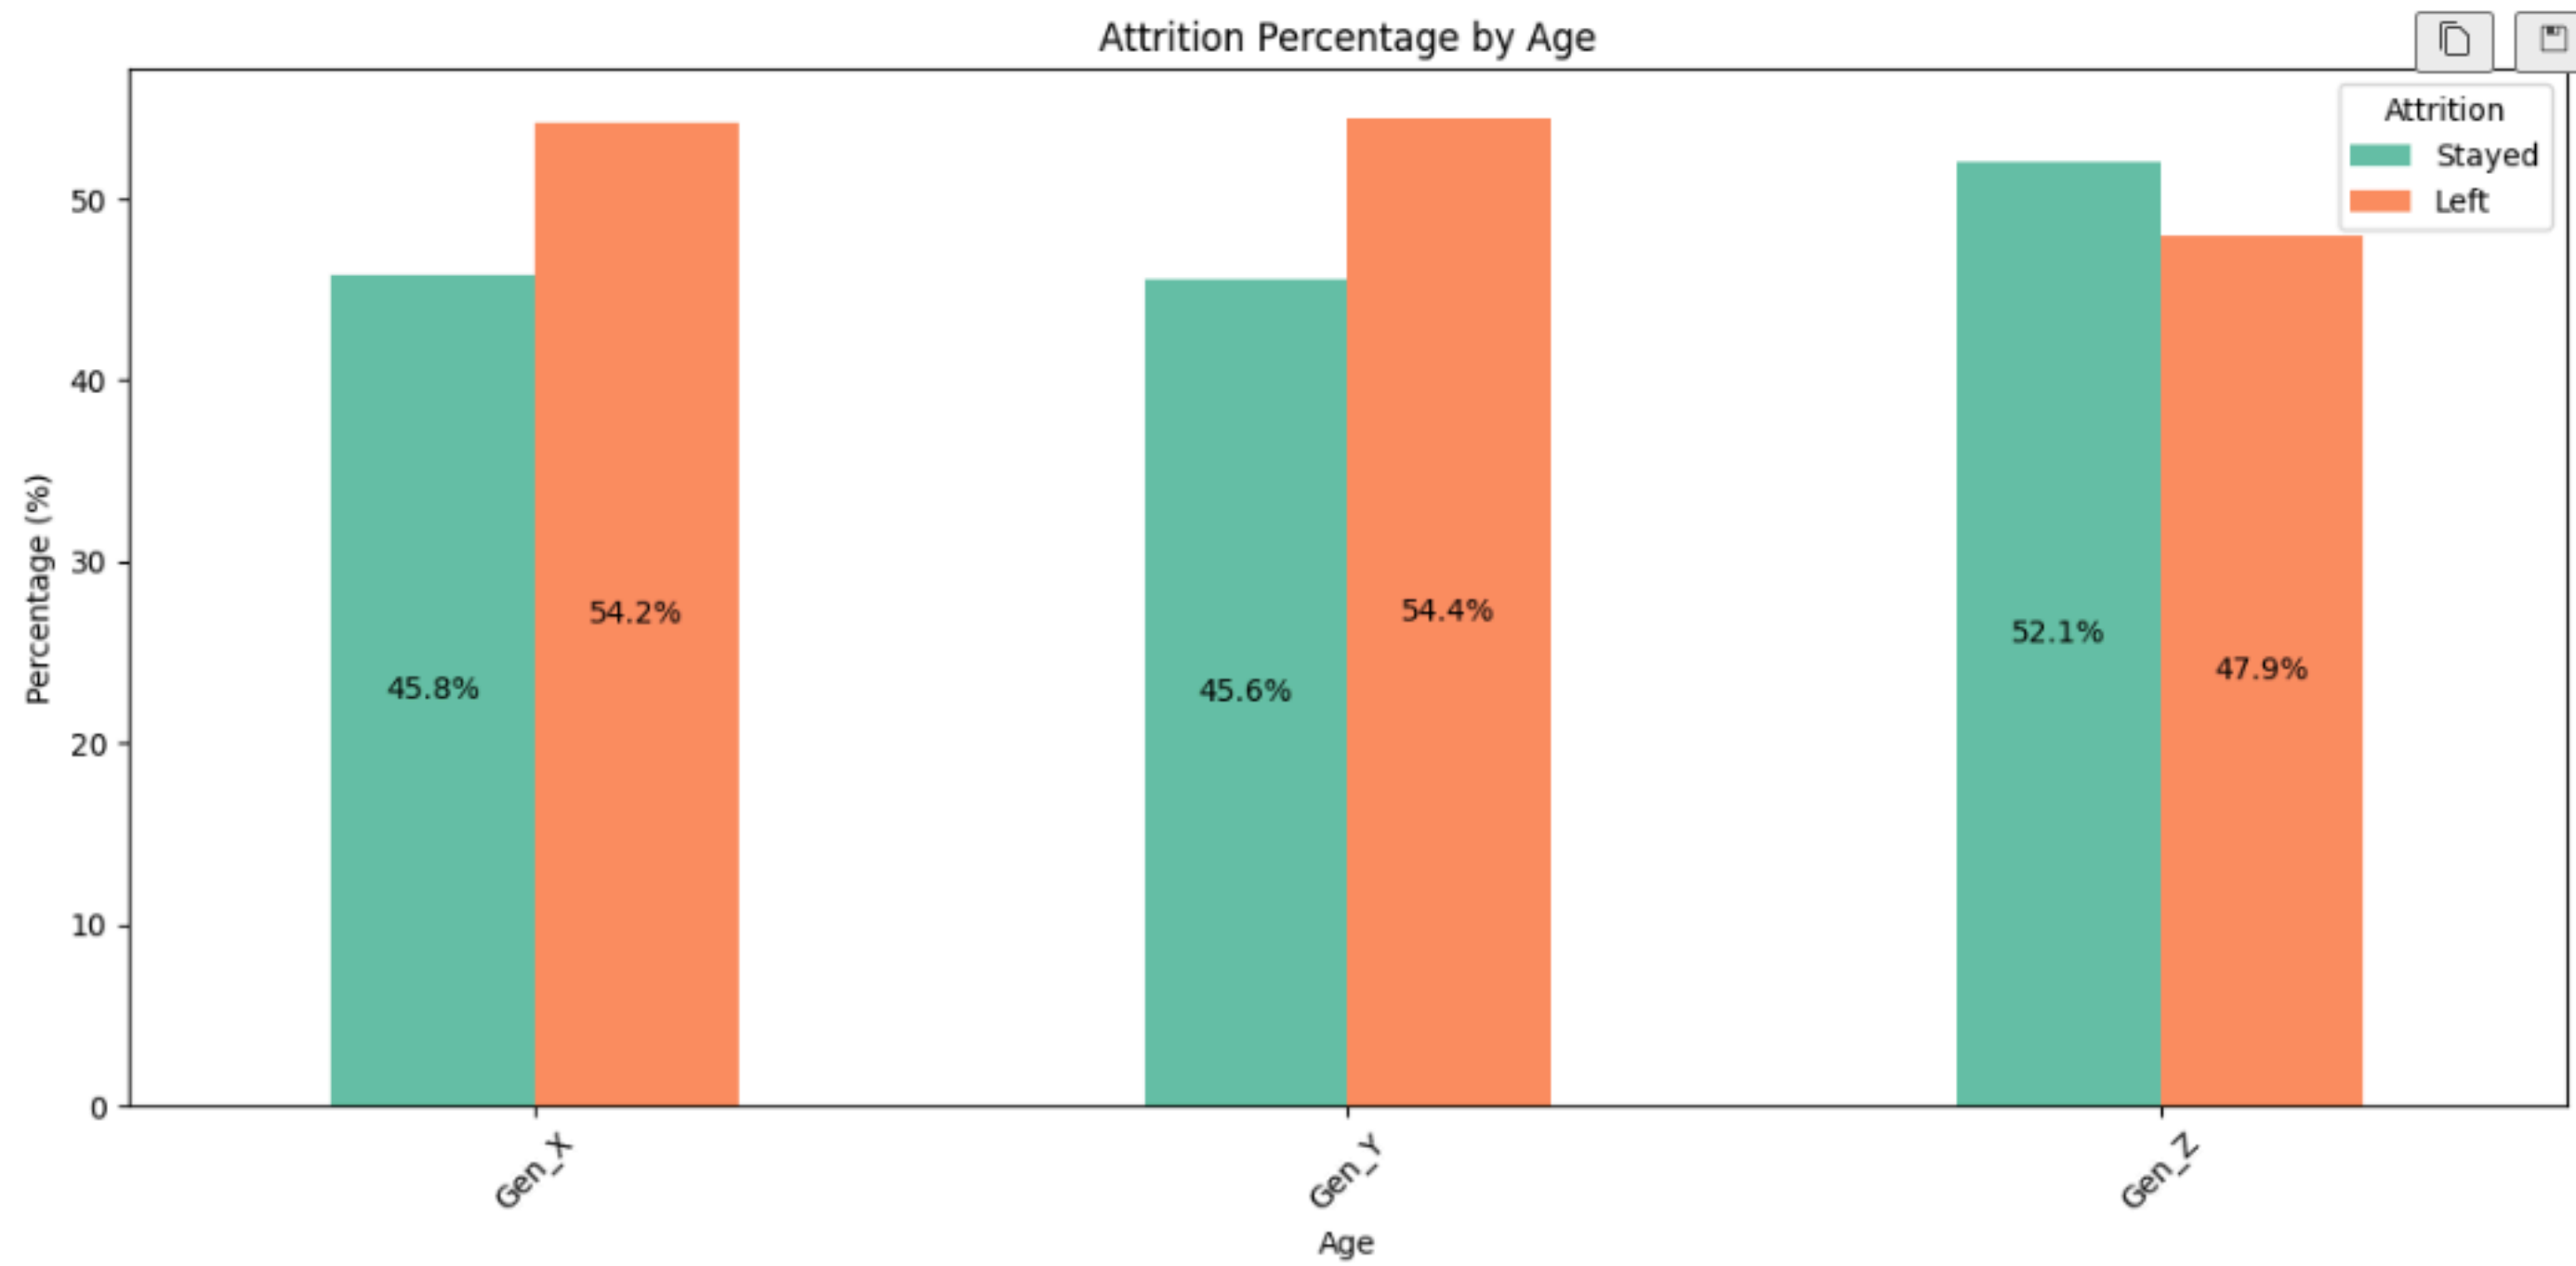
✓ 0.1s

Years at

Monthly

Work-

Id



Data Preparation

```
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder
continuous_cols = ['Age', 'Years at Company', 'Monthly Income', 'Number of Promotions', 'Distance from Home', 'Company Tenure']
scaler = StandardScaler()
data[continuous_cols] = scaler.fit_transform(data[continuous_cols])
```

Python

```
(count      7.449800e+04
mean       2.114517e-16
std        1.000007e+00
min        -1.699008e+00
25%        -8.714242e-01
50%         3.891746e-02
75%         8.665008e-01
max         1.694084e+00
Name: Age, dtype: float64,
Age
-0.374874    1875
 0.121676    1861
 1.197534    1843
 0.535467    1842
-0.043841    1834
 1.363051    1824
 0.369951    1822
-0.705908    1818
 0.038917    1813
 0.700984    1806
-0.292116    1806
-0.209358    1795
 0.618226    1793
-0.457633    1791
-1.616249    1789
...
-1.285216    1711
 1.114776    1708
 0.783742    1702
-1.699008    1702
```



```
label_encoder = LabelEncoder()  
data['Attrition'] = label_encoder.fit_transform(data['Attrition'])
```

```
data['Attrition'].value_counts()
```

```
Attrition  
1    39128  
0    35370  
Name: count, dtype: int64
```

1 คือ อยู่ต่อ
0 คือ ลาออก

```
# คอลัมน์ที่เป็นตัวแปรประเภท categorical ที่ต้องการทำ One-Hot Encoding
categorical_cols = ['Gender', 'Job Role', 'Work-Life Balance', 'Job Satisfaction',
                    'Performance Rating', 'Marital Status', 'Education Level',
                    'Job Level', 'Company Size', 'Remote Work',
                    'Leadership Opportunities', 'Innovation Opportunities',
                    'Company Reputation', 'Employee Recognition', 'Overtime']


# ทำ One-Hot Encoding โดยใช้ pd.get_dummies()
data = pd.get_dummies(data, columns=categorical_cols)

# แสดงผลลัพธ์
print(data.head())
```


57

```
Index(['Age', 'Years at Company', 'Monthly Income', 'Number of Promotions',  
      'Distance from Home', 'Number of Dependents', 'Company Tenure',  
      'Attrition', 'Gender_Female', 'Gender_Male', 'Job Role_Education',  
      'Job Role_Finance', 'Job Role_Healthcare', 'Job Role_Media',  
      'Job Role_Technology', 'Work-Life Balance_Excellent',  
      'Work-Life Balance_Fair', 'Work-Life Balance_Good',  
      'Work-Life Balance_Poor', 'Job Satisfaction_High',  
      'Job Satisfaction_Low', 'Job Satisfaction_Medium',  
      'Job Satisfaction_Very High', 'Performance Rating_Average',  
      'Performance Rating_Below Average', 'Performance Rating_High',  
      'Performance Rating_Low', 'Marital Status_Divorced',  
      'Marital Status_Married', 'Marital Status_Single',  
      'Education Level_Associate Degree', 'Education Level_Bachelor's Degree',  
      'Education Level_High School', 'Education Level_Master's Degree',  
      'Education Level_PhD', 'Job Level_Entry', 'Job Level_Mid',  
      'Job Level_Senior', 'Company Size_Large', 'Company Size_Medium',  
      'Company Size_Small', 'Remote Work_No', 'Remote Work_Yes',  
      'Leadership Opportunities_No', 'Leadership Opportunities_Yes',  
      'Innovation Opportunities_No', 'Innovation Opportunities_Yes',  
      'Company Reputation_Excellent', 'Company Reputation_Fair',  
      'Company Reputation_Good', 'Company Reputation_Poor',  
      'Employee Recognition_High', 'Employee Recognition_Low',  
      'Employee Recognition_Medium', 'Employee Recognition_Very High',  
      'Overtime_No', 'Overtime_Yes'],  
      dtype='object')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74498 entries, 0 to 74497
Data columns (total 57 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   74498 non-null  float64
1   Years at Company                     74498 non-null  float64
2   Monthly Income                       74498 non-null  float64
3   Number of Promotions                 74498 non-null  float64
4   Distance from Home                   74498 non-null  float64
5   Number of Dependents                 74498 non-null  int64
6   Company Tenure                       74498 non-null  float64
7   Attrition                           74498 non-null  int64
8   Gender_Female                       74498 non-null  bool
9   Gender_Male                         74498 non-null  bool
10  Job_Role_Education                  74498 non-null  bool
11  Job_Role_Finance                    74498 non-null  bool
12  Job_Role_Healthcare                 74498 non-null  bool
13  Job_Role_Media                      74498 non-null  bool
14  Job_Role_Technology                 74498 non-null  bool
15  Work-Life Balance_Excellent          74498 non-null  bool
16  Work-Life Balance_Fair               74498 non-null  bool
17  Work-Life Balance_Good               74498 non-null  bool
18  Work-Life Balance_Poor               74498 non-null  bool
19  Job_Satisfaction_High                74498 non-null  bool
...
56  Overtime_Yes                        74498 non-null  bool
dtypes: bool(49), float64(6), int64(2)
```



ถ้า Correalation Matrix มีค่า 0.03 ขึ้นไป
แสดงค่านั้นมีผลต่อการลาออกมาก



```
data = data.drop(['Gender_Female', 'Gender_Male', 'Job Role_Education', 'Job Role_Finance', 'Job Role_Healthcare', 'Job Role_Media',  
|               'Job Role_Technology']), axis=1)  
print(data.columns)
```

Python

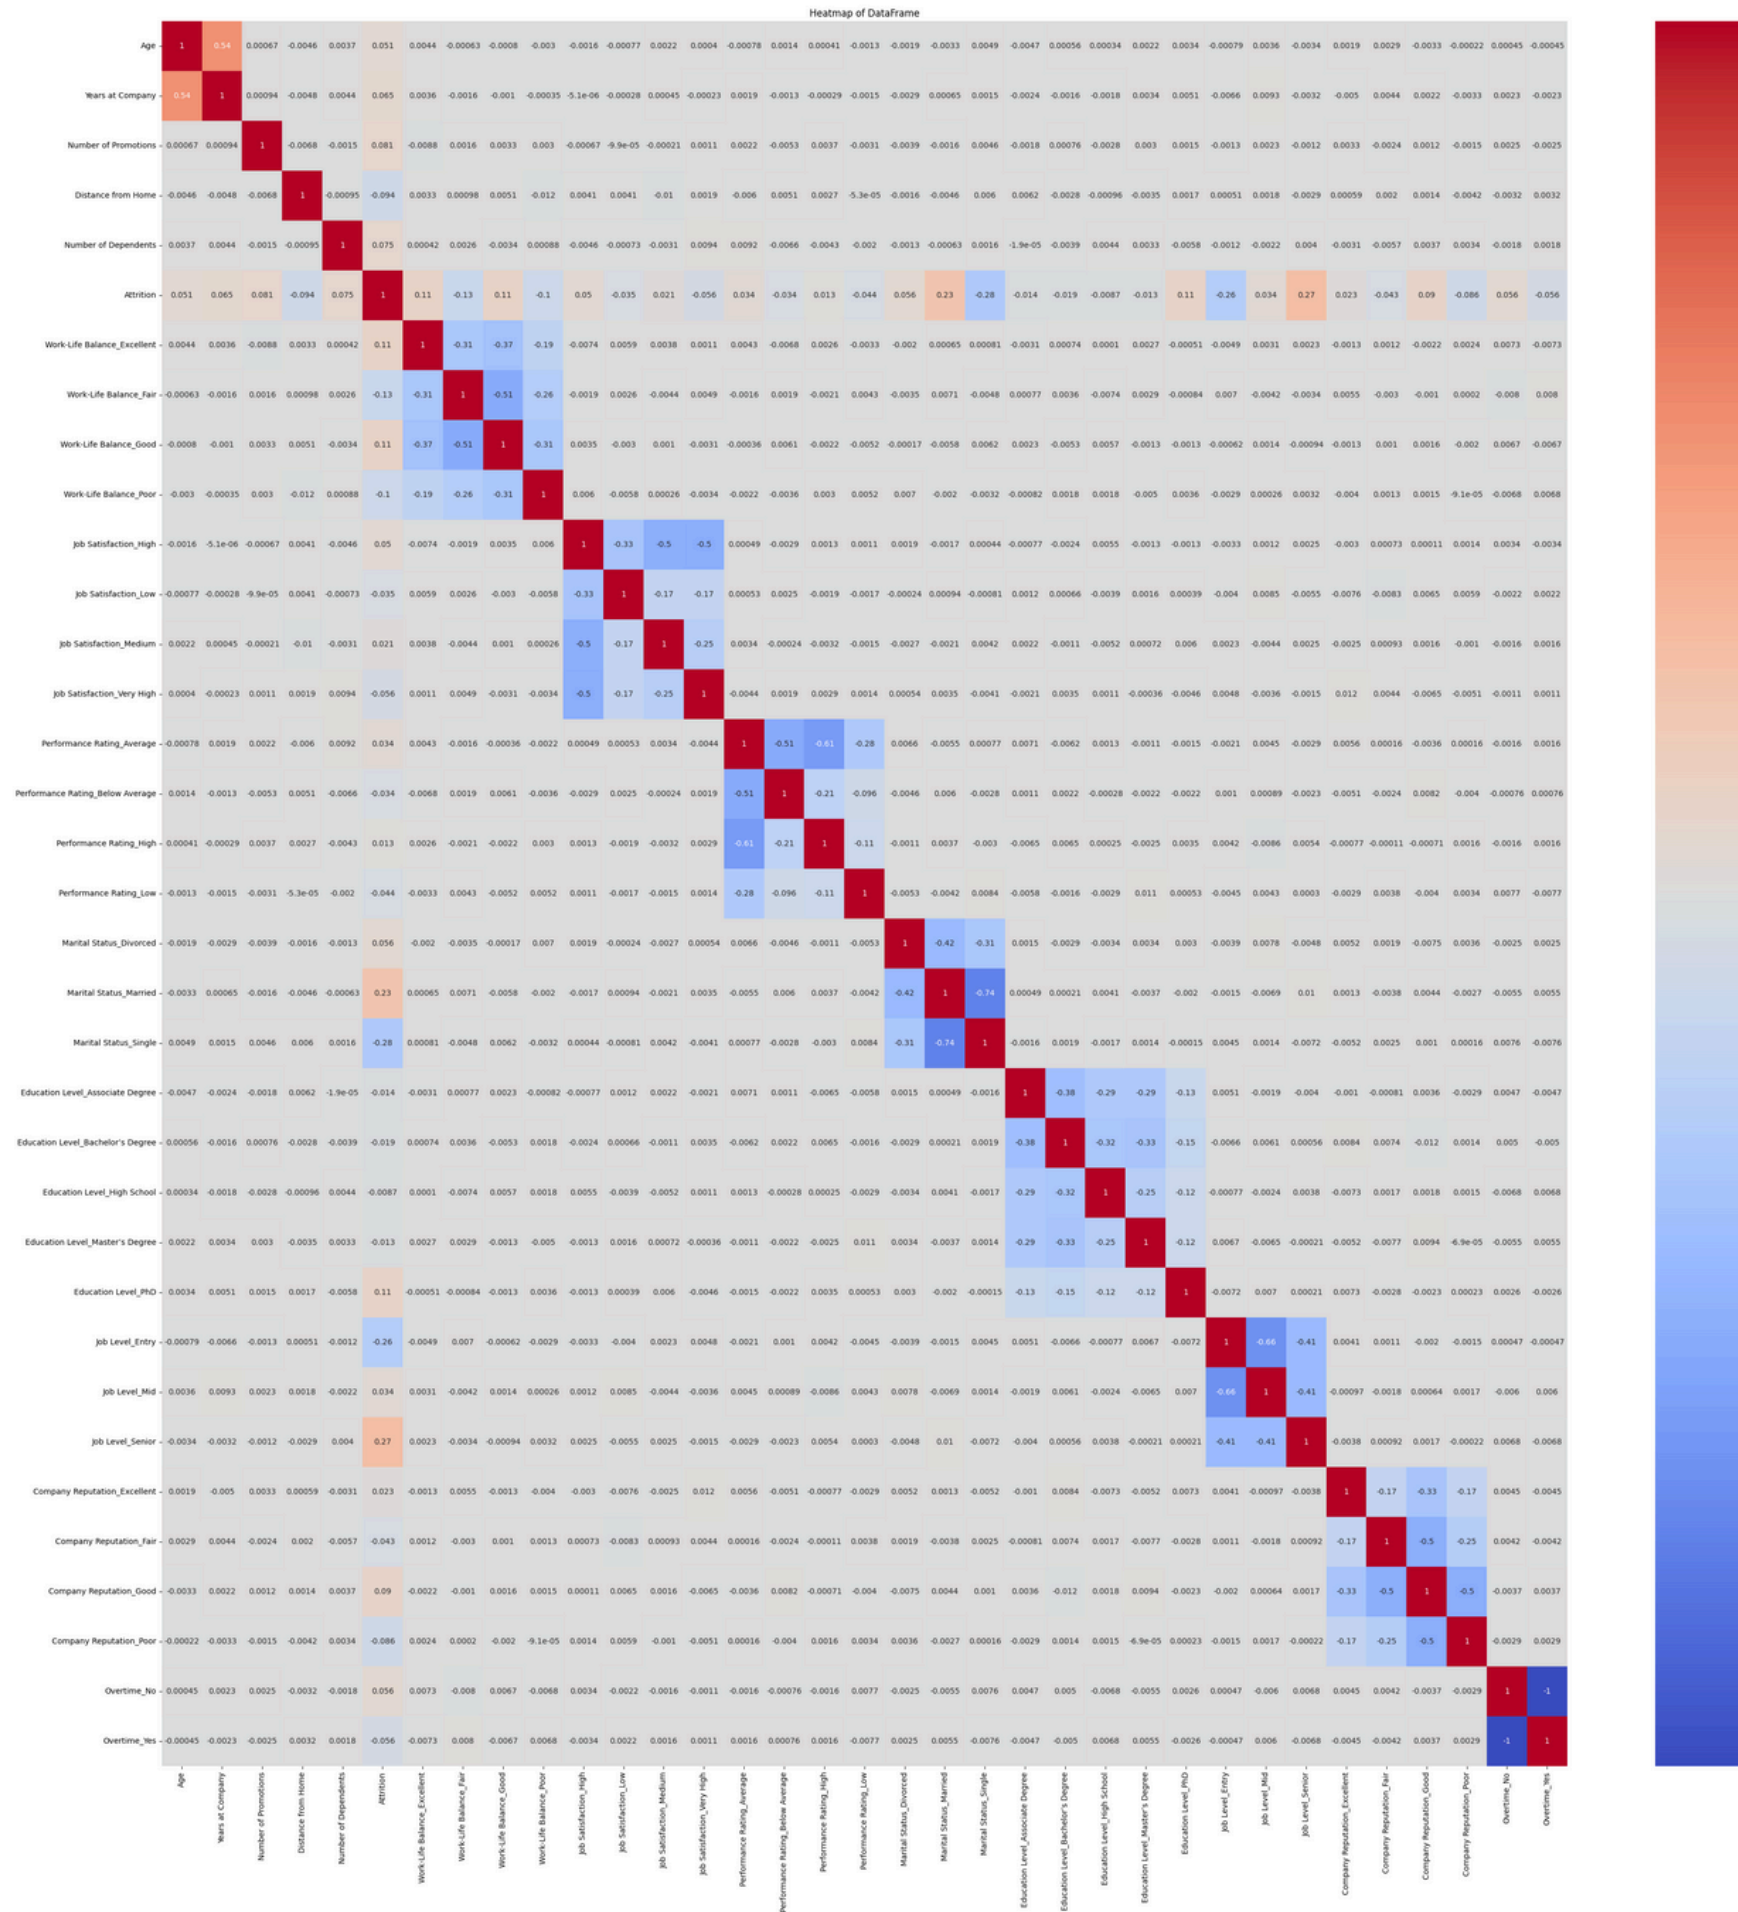
```
data = data.drop(['Employee_Recognition_Medium', 'Employee_Recognition_Low', 'Employee_Recognition_High', 'Employee_Recognition_Very  
|           'Innovation_Opportunities_Yes', 'Leadership_Opportunities_No', 'Leadership_Opportunities_Yes', 'Company_Size_Medium', 'Company_S  
|           'Remote_Work_No', 'Remote_Work_Yes', 'Company_Tenure', 'Monthly_Income', 'Company_Size_Large']), axis=1)
```

Python


```
print(len(data.columns))  
print(data.columns)
```

35

```
Index(['Age', 'Years at Company', 'Number of Promotions', 'Distance from Home',  
      'Number of Dependents', 'Attrition', 'Work-Life Balance_Excellent',  
      'Work-Life Balance_Fair', 'Work-Life Balance_Good',  
      'Work-Life Balance_Poor', 'Job Satisfaction_High',  
      'Job Satisfaction_Low', 'Job Satisfaction_Medium',  
      'Job Satisfaction_Very High', 'Performance Rating_Average',  
      'Performance Rating_Below Average', 'Performance Rating_High',  
      'Performance Rating_Low', 'Marital Status_Divorced',  
      'Marital Status_Married', 'Marital Status_Single',  
      'Education Level_Associate Degree', 'Education Level_Bachelor's Degree',  
      'Education Level_High School', 'Education Level_Master's Degree',  
      'Education Level_PhD', 'Job Level_Entry', 'Job Level_Mid',  
      'Job Level_Senior', 'Company Reputation_Excellent',  
      'Company Reputation_Fair', 'Company Reputation_Good',  
      'Company Reputation_Poor', 'Overtime_No', 'Overtime_Yes'],  
      dtype='object')
```

```
sns.countplot(x='Attrition', data=data)  
plt.show()
```



```
# แบ่งข้อมูล
X = data.drop(columns=['Attrition']) # Features (ทั้งหมดยกเว้น Attrition)
y = data['Attrition']                # Target column

# แบ่งข้อมูล
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print(f"Train set size: {X_train.shape}, Test set size: {X_test.shape}")
```

```
Train set size: (59598, 34), Test set size: (14900, 34)
```

```

model = LogisticRegression(random_state=42)

# ฝึกโมเดล
model.fit(X_train, y_train)

# ทดสอบโมเดล
y_pred = model.predict(X_test)

# ประเมินโมเดล
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

cftmt_lo = confusion_matrix(y_test, y_pred)
print(cftmt_lo)
custom_plt_confution_matrix(cftmt_lo)


```

```

Accuracy: 0.7322147651006712

```

	precision	recall	f1-score	support
0	0.72	0.71	0.72	7096
1	0.74	0.76	0.75	7804
accuracy			0.73	14900
macro avg	0.73	0.73	0.73	14900
weighted avg	0.73	0.73	0.73	14900



```

gb_model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
max_depth=3, random_state=42)
# เทรนโมเดลด้วยข้อมูล train
gb_model.fit(X_train, y_train)

# ทำนายผลด้วยข้อมูล test
y_pred_gb = gb_model.predict(X_test)

# ประเมินผลโมเดล
print("Accuracy:", accuracy_score(y_test, y_pred_gb))
print("\nClassification Report:\n", classification_report(y_test, y_pred_gb))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_gb))

gBoostCunMatrix = confusion_matrix(y_test, y_pred_gb)


```

Accuracy: 0.7367785234899329

Classification Report:

	precision	recall	f1-score	support
0	0.73	0.71	0.72	7096
1	0.74	0.76	0.75	7804
accuracy			0.74	14900
macro avg	0.74	0.74	0.74	14900
weighted avg	0.74	0.74	0.74	14900

Seaborn Confusion Matrix with labels



```
k = 5


# สร้างโมเดล Gradient Boosting
gb_model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)

# สร้าง K-Fold Cross-Validator (Stratified เพื่อให้คลาสสมดุลในแต่ละ Fold)
kfold = StratifiedKFold(n_splits=k, shuffle=True, random_state=42)

# ใช้ cross_val_score ประเมินโมเดล
scores = cross_val_score(gb_model, X, y, cv=kfold, scoring='accuracy')

# แสดงผลคะแนน
print(f"K-Fold Cross-Validation Results ({k} folds):")
print(f"Scores: {scores}")
print(f"Mean Accuracy: {scores.mean():.2f}")
print(f"Standard Deviation: {scores.std():.2f}")
```

```
K-Fold Cross-Validation Results (5 folds):
Scores: [0.72865772 0.72516779 0.73328859 0.74045238 0.73380764]
Mean Accuracy: 0.73
Standard Deviation: 0.01
```



```

model = LogisticRegression(random_state=42)

# ฝึกโมเดล
model.fit(X_train, y_train)

# ทดสอบโมเดล
y_pred = model.predict(X_test)

# ประเมินโมเดล
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
cftmt_lo = confusion_matrix(y_test, y_pred)
print(cftmt_lo)
custom_plt_confution_matrix(cftmt_lo)

```


```

Accuracy: 0.5828187919463087

```

	precision	recall	f1-score	support
0	0.57	0.53	0.55	7096
1	0.60	0.63	0.61	7804
accuracy			0.58	14900
macro avg	0.58	0.58	0.58	14900
weighted avg	0.58	0.58	0.58	14900

Seaborn Confusion Matrix with labels




```

gb_model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, max_depth=3,
random_state=42)
# เทรนโมเดลด้วยข้อมูล train
gb_model.fit(X_train, y_train)

# ทำนายผลด้วยข้อมูล test
y_pred_gb = gb_model.predict(X_test)

# ประเมินผลโมเดล
print("Accuracy:", accuracy_score(y_test, y_pred_gb))
print("\nClassification Report:\n", classification_report(y_test, y_pred_gb))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_gb))
gBoostConMatrix = confusion_matrix(y_test, y_pred_gb)

```

Accuracy: 0.7622147651006711

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.74	0.75	7096
1	0.77	0.78	0.77	7804
accuracy			0.76	14900
macro avg	0.76	0.76	0.76	14900
weighted avg	0.76	0.76	0.76	14900

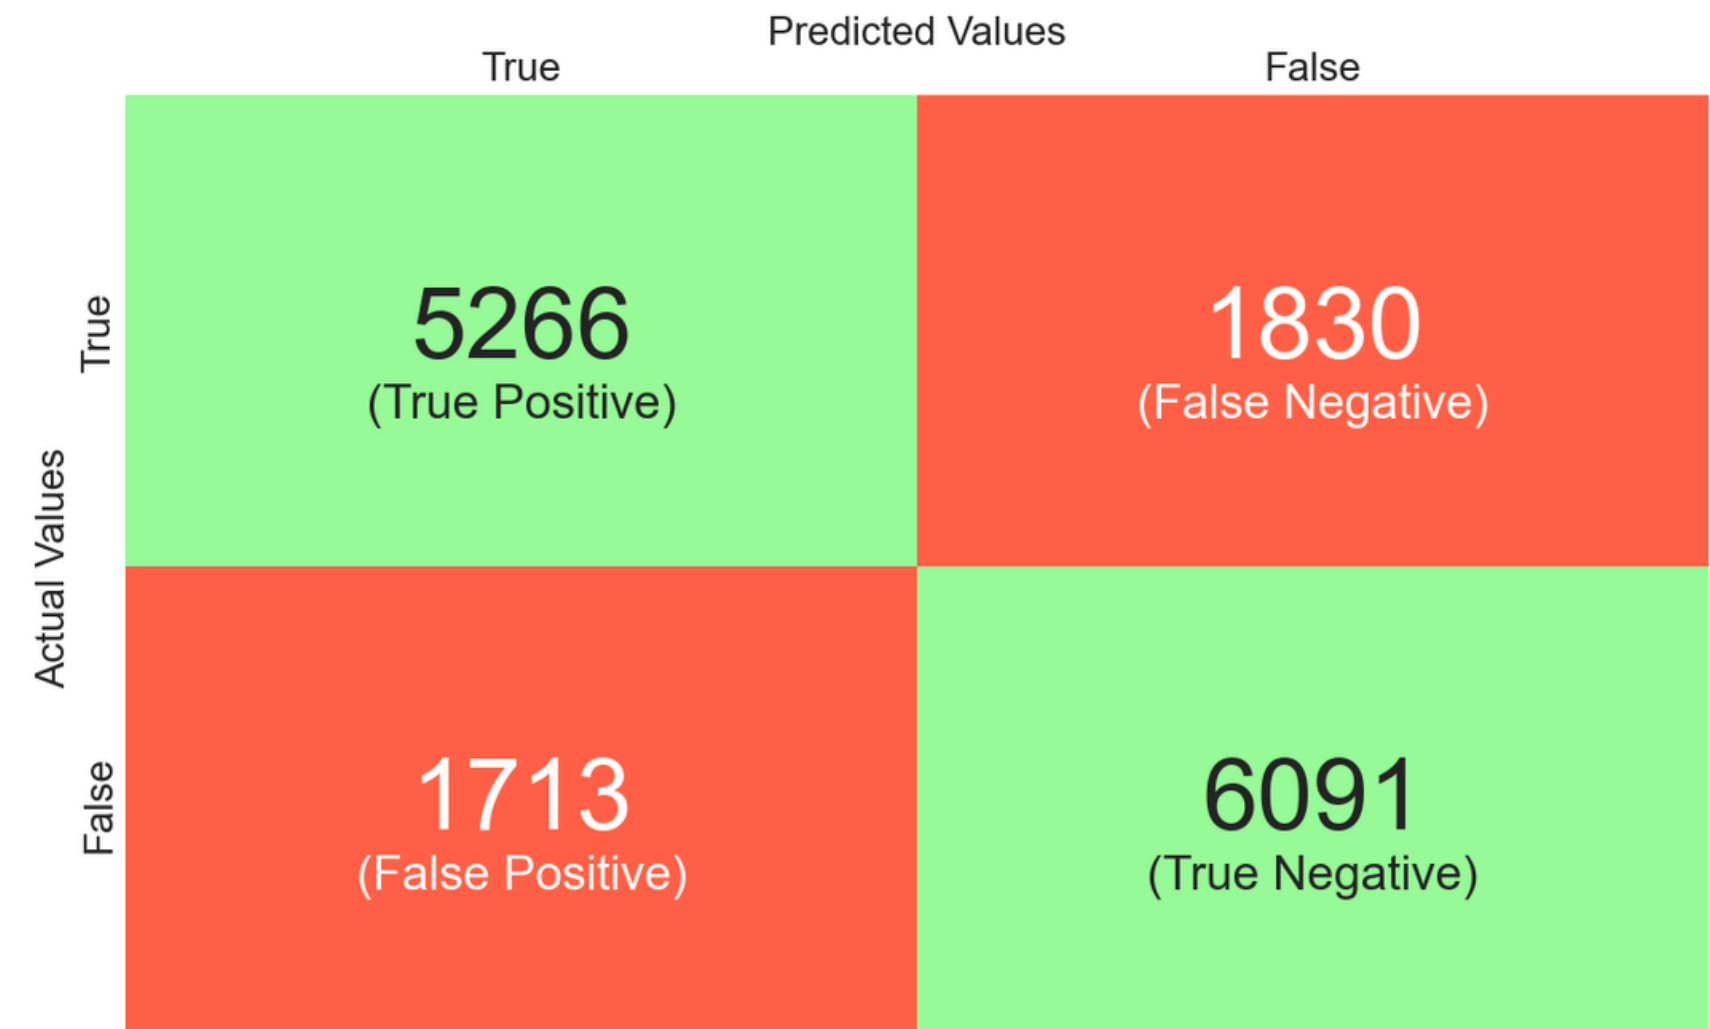
Confusion Matrix:

```

[[5266 1830]
 [1713 6091]]

```

Seaborn Confusion Matrix with labels



Logistic Regression

เปรียบเทียบความแตกต่าง

เปรียบเทียบ	แบบไม่ทำอะไรเลย	แบบทำ Standardization & Drop features
Accuracy	58.2%	73.2%
Precision (0,1)	(0.57, 0.60)	(0.72, 0.74)
Recall (0,1)	(0.53, 0.63)	(0.71, 0.76)
Confusion Matrix (TP, FN, FP, TN)	(3773, 3323, 2893, 4911)	(5008, 2088, 1902, 5902)

Gradient Boosting

เปรียบเทียบความแตกต่าง

Metric	Gradient Boosting (ไม่ทำอะไรเลย)	Gradient Boosting (ทำ Feature Engineering)
Accuracy	76.2%	73.6%
Precision (0,1)	(0.75, 0.77)	(0.73, 0.74)
Recall (0,1)	(0.75, 0.78)	(0.71, 0.76)
Confusion Matrix (TP, FN, FP, TN)	(5266, 1830, 1713, 6091)	(5058, 2038, 1884, 5920)

```
data = data.drop(['Marital Status_Divorced', 'Marital Status_Married',  
                  'Marital Status_Single']), axis=1)  
print(data.columns)
```

```
gb_model.fit(X_train, y_train)  
  
# ทำนายผลด้วยข้อมูล test  
y_pred_gb = gb_model.predict(X_test)  
  
# ประเมินผลโมเดล  
print("Accuracy:", accuracy_score(y_test, y_pred_gb))  
print("\nClassification Report:\n", classification_report(y_test, y_pred_gb))  
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_gb))  
  
gBoostCunMatrix = confusion_matrix(y_test, y_pred_gb)
```

[42] ✓ 7.3s Python

... Accuracy: 0.6989261744966443

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.69	0.69	7096
1	0.72	0.71	0.71	7804
accuracy			0.70	14900
macro avg	0.70	0.70	0.70	14900
weighted avg	0.70	0.70	0.70	14900

Confusion Matrix:

```
[[4907 2189]  
 [2297 5507]]
```

Logistic Regression

```
Accuracy: 0.6989261744966443
```

Gradient Boosting

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.69	0.69	7096
1	0.72	0.71	0.71	7804
accuracy			0.70	14900
macro avg	0.70	0.70	0.70	14900
weighted avg	0.70	0.70	0.70	14900

Confusion Matrix:

```
[[4907 2189]  
 [2297 5507]]
```

หลังจากถอดรอบ Feature ที่มีผลต่อ y มากที่สุดออก

มีค่า Accuracy: 0.6 น้อยลง

```
data = data.drop(['Gender_Female', 'Gender_Male' , 'Job Role_Education' ,
'Job Role_Finance', 'Job Role_Healthcare', 'Job Role_Media',
'Job Role_Technology']), axis=1)
print(data.columns)
```

Accuracy: 0.7331543624161074

	precision	recall	f1-score	support
0	0.73	0.71	0.72	7096
1	0.74	0.76	0.75	7804
accuracy			0.73	14900
macro avg	0.73	0.73	0.73	14900
weighted avg	0.73	0.73	0.73	14900

[[5009 2087]
[1889 5915]]

Logistic Regression

Accuracy: 0.7377852348993289

Classification Report:

	precision	recall	f1-score	support
0	0.73	0.71	0.72	7096
1	0.74	0.76	0.75	7804
accuracy			0.74	14900
macro avg	0.74	0.74	0.74	14900
weighted avg	0.74	0.74	0.74	14900

Confusion Matrix:

[[5040 2056]
[1851 5953]]

Gradient Boosting

หลังจากลบ Feature ที่มีค่า y น้อย
มีค่า Accuracy: 0.73 ไม่ต่างกันมาก

Note

1. ถ้า Correalation Matrix มีค่า 0.03 ขึ้นไปให้ใช้ค่านั้น (ว่ามีผลต่อการลาออก)
2. [1 คืออยู่, 0 คือลาออก]
3. (ตอนนี้เรา predict ว่าปัจจัยใดที่มีผลต่อการลาออกได้แล้ว) [27 Jan 2025]